

Multiple Linear Regression Case Study: Boston Housing Prices

Background: The data has 506 cases where each case is a location in Boston. The “median housing price” is a target variable. The data has many other variables related to environment, education, crime etc. which can influence the housing prices in the specific location

The objective is to identify significant factors affecting housing prices.

Import data and display first 6 rows

```
import os
import pandas as pd

os.chdir("C:/Sankhya")
boston = pd.read_csv('Housing Prices.csv')
boston.head(6)
```

	CRIM	ZN	INDUS	CHAS	NOX	...	RAD	TAX	PTRATIO	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	...	1	296	15.3	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	...	2	242	17.8	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	...	2	242	17.8	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	...	3	222	18.7	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	...	3	222	18.7	5.33	36.2
5	0.02985	0.0	2.18	0	0.458	...	3	222	18.7	5.21	28.7

[6 rows x 13 columns]

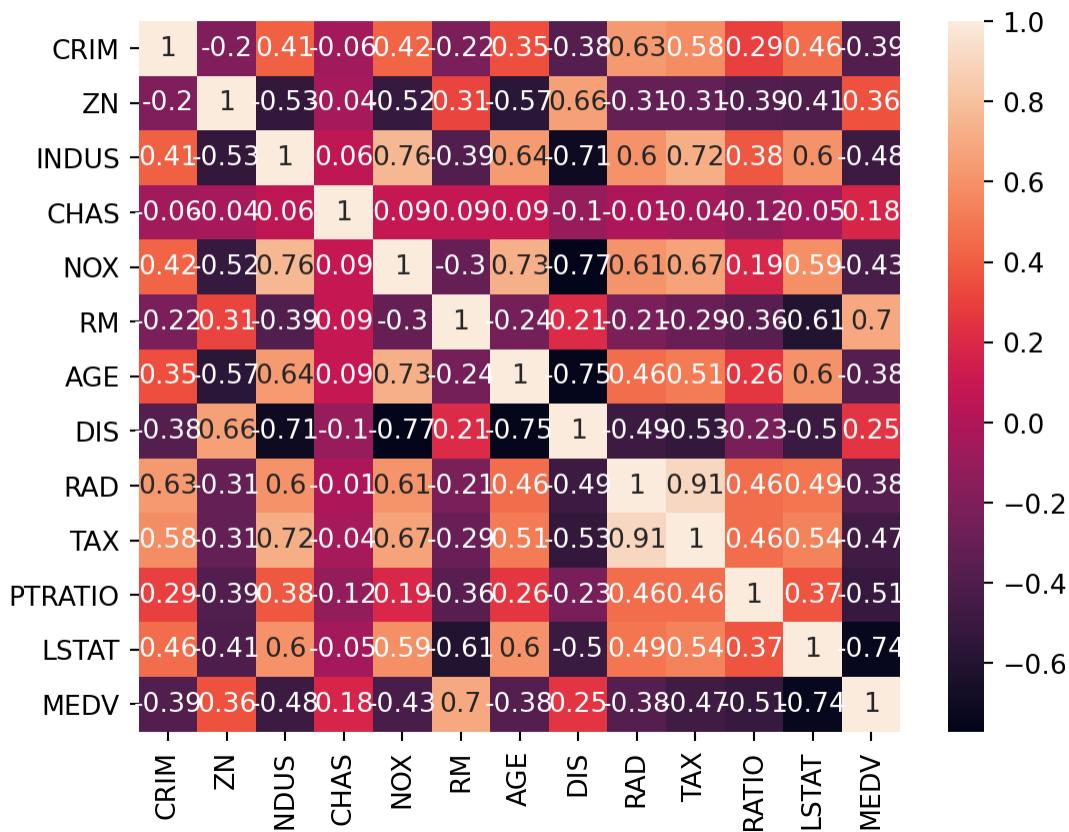
Data Description

Column name	Column description
CRIM	per capita crime rate by town
ZN	proportion of residential land zoned for lots over 25,000 in sq ft
INDUS	proportion of non-retail business acres per town %
CHAS	Charles River dummy variable (=1 if tract bounds river; 0 otherwise)

Column name	Column description
NOX	nitric oxides concentration (parts per 10 million)
RM	average number of rooms per dwelling
AGE	proportion of owner-occupied units built prior to 1940 %
DIS	weighted distances to five Boston employment centres
RAD	index of accessibility to radial highways
TAX	full-value property-tax rate per \$10,000
PTRATIO	pupil-teacher ratio by town
LSAT	% lower status of the population
MEDV	Median value of owner-occupied homes in \$1000s

Correlation matrix using heatmap

```
import seaborn as sns
correlation_matrix = boston.corr().round(2)
sns.heatmap(data=correlation_matrix, annot=True)
```



Check for multicollinearity using vif

```

from patsy import dmatrices
from statsmodels.stats.outliers_influence import variance_inflation_factor

y, X = dmatrices('MEDV~CRIM+ZN+INDUS+CHAS+NOX+RM+AGE+DIS+RAD+TAX+PTRATIO+LSTAT', data=boston,
return_type="dataframe")

vif = pd.Series([variance_inflation_factor(X.values, i)for i in range(X.shape[1])],index=X.co
lumns)
vif

```

```

Intercept      535.526619
CRIM           1.767486
ZN              2.298459
INDUS          3.987181
CHAS            1.071168
NOX             4.369093
RM              1.912532
AGE             3.088232
DIS             3.954037
RAD             7.445301
TAX             9.002158
PTRATIO         1.797060
LSTAT           2.870777
dtype: float64

```

Comment: It is observed that variable TAX has high vif

The variable “TAX” is excluded

```
boston = boston.drop(['TAX'],1)
```

Check for multicollinearity again

```

y2, X2 = dmatrices('MEDV~CRIM+ZN+INDUS+CHAS+NOX+RM+AGE+DIS+RAD+PTRATIO+LSTAT', data=boston,
return_type="dataframe")

vif = pd.Series([variance_inflation_factor(X2.values, i)for i in range(X2.shape[1])],index=X
2.columns)
vif

```

```
Intercept      529.480235
CRIM           1.767349
ZN              2.184172
INDUS          3.217951
CHAS            1.055023
NOX             4.343300
RM              1.902642
AGE             3.085756
DIS             3.952445
RAD              2.772208
PTRATIO         1.787049
LSTAT            2.870408
dtype: float64
```

The multicollinearity problem is resolved as all VIF's are less than 5

Develop Housing Prices Model using Linear Regression

Split original data into training and testing data sets

```
from sklearn.model_selection import train_test_split
boston_train, boston_test = train_test_split(boston, test_size=0.2, random_state=42)
```

Running Linear Regression model using OLS function from statsmodels library

```
import statsmodels.formula.api as smf

hp_model = smf.ols('MEDV~CRIM+ZN+INDUS+CHAS+NOX+RM+AGE+DIS+RAD+PTRATIO+LSTAT', data=boston_train).fit()
```

Displaying model summary

```
print(hp_model.summary())
```

OLS Regression Results

Dep. Variable:	MEDV	R-squared:	0.735			
Model:	OLS	Adj. R-squared:	0.728			
Method:	Least Squares	F-statistic:	98.99			
Date:	Tue, 19 Dec 2023	Prob (F-statistic):	7.75e-106			
Time:	10:55:50	Log-Likelihood:	-1206.6			
No. Observations:	404	AIC:	2437.			
Df Residuals:	392	BIC:	2485.			
Df Model:	11					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	35.2203	5.555	6.340	0.000	24.299	46.141
CRIM	-0.1266	0.036	-3.556	0.000	-0.197	-0.057
ZN	0.0189	0.016	1.196	0.232	-0.012	0.050
INDUS	-0.0452	0.064	-0.712	0.477	-0.170	0.080
CHAS	3.3202	0.975	3.406	0.001	1.404	5.237
NOX	-19.3122	4.327	-4.463	0.000	-27.819	-10.805
RM	4.3369	0.473	9.168	0.000	3.407	5.267
AGE	-0.0034	0.015	-0.225	0.822	-0.033	0.026
DIS	-1.4400	0.232	-6.205	0.000	-1.896	-0.984
RAD	0.0794	0.047	1.687	0.092	-0.013	0.172
PTRATIO	-0.9248	0.148	-6.269	0.000	-1.215	-0.635
LSTAT	-0.5345	0.057	-9.412	0.000	-0.646	-0.423
Omnibus:	122.472	Durbin-Watson:	2.171			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	473.320			
Skew:	1.302	Prob(JB):	1.66e-103			
Kurtosis:	7.620	Cond. No.	2.06e+03			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.06e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Comment: From the model summary, it can be seen that the variables ZN, INDUS,RAD and AGE are insignificant(since p value > 0.05), so we remove them from the model and run the model again.

```
hp_model = smf.ols('MEDV~CRIM+CHAS+NOX+RM+DIS+PTRATIO+LSTAT', data=boston_train).fit()
print(hp_model.summary())
```

OLS Regression Results

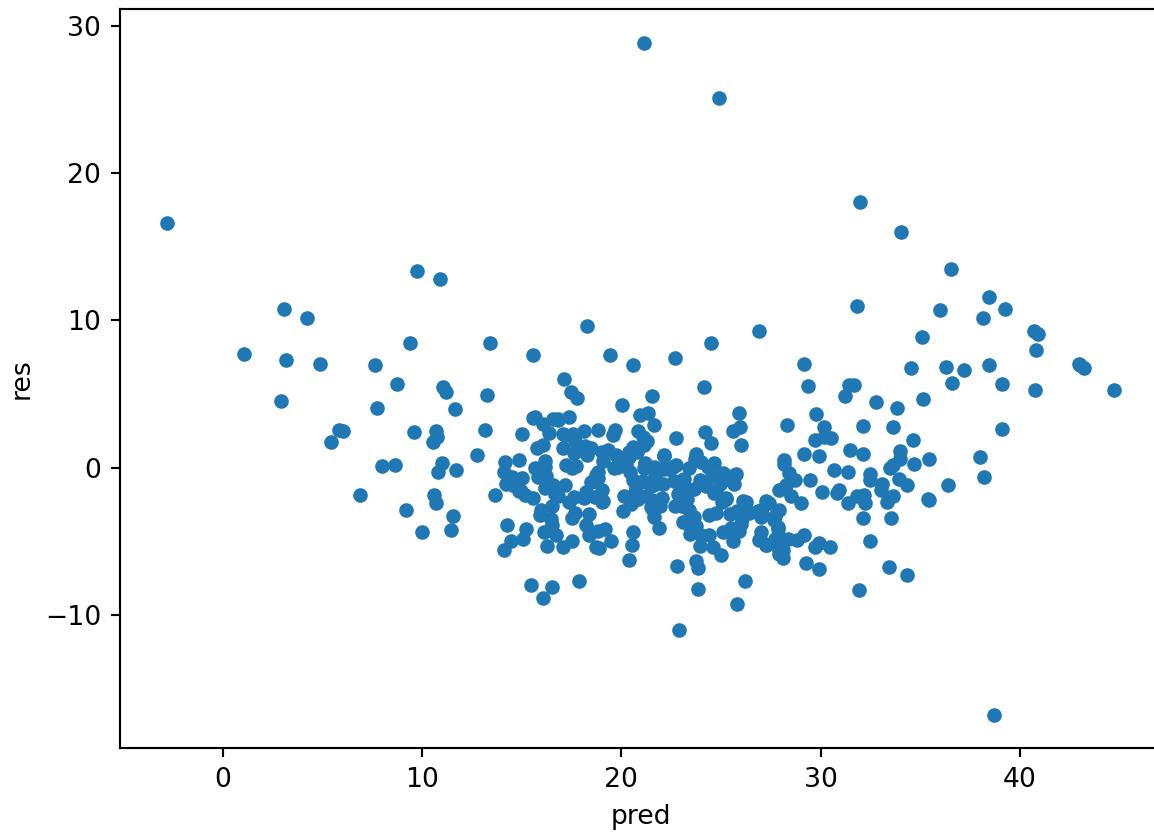
Dep. Variable:	MEDV	R-squared:	0.732			
Model:	OLS	Adj. R-squared:	0.727			
Method:	Least Squares	F-statistic:	154.2			
Date:	Tue, 19 Dec 2023	Prob (F-statistic):	6.26e-109			
Time:	10:55:51	Log-Likelihood:	-1209.4			
No. Observations:	404	AIC:	2435.			
Df Residuals:	396	BIC:	2467.			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	31.8452	5.073	6.278	0.000	21.872	41.818
CRIM	-0.0925	0.031	-2.938	0.003	-0.154	-0.031
CHAS	3.3351	0.974	3.425	0.001	1.421	5.249
NOX	-18.2416	3.613	-5.049	0.000	-25.344	-11.139
RM	4.5754	0.447	10.225	0.000	3.696	5.455
DIS	-1.2243	0.184	-6.660	0.000	-1.586	-0.863
PTRATIO	-0.8962	0.122	-7.317	0.000	-1.137	-0.655
LSTAT	-0.5316	0.054	-9.842	0.000	-0.638	-0.425
Omnibus:	137.650	Durbin-Watson:	2.160			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	610.512			
Skew:	1.426	Prob(JB):	2.69e-133			
Kurtosis:	8.304	Cond. No.	558.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot of Residuals vs Predicted values

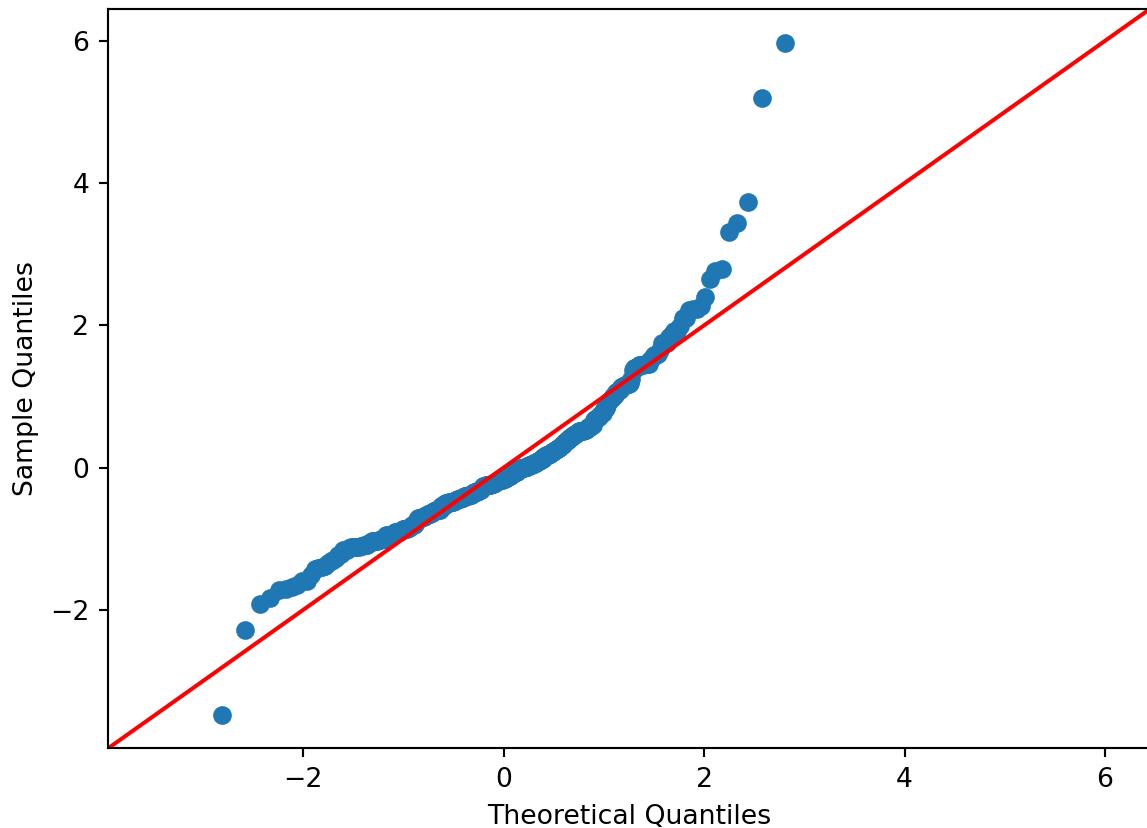
```
boston_train['pred'] = hp_model.predict(boston_train)
boston_train['res']=hp_model.resid
boston_train.plot.scatter(x='pred', y='res')
```



Comment: It is observed that residuals are randomly distributed and uncorrelated with predicted values

Check if distribution of errors is “NORMAL”

```
import statsmodels.api as sm
sm.graphics.qqplot(boston_train.res, line='45', fit=True)
```



```
import scipy as sp
sp.stats.shapiro(boston_train.res)
```

```
ShapiroResult(statistic=0.9116807579994202, pvalue=1.2912867356060145e-14)
```

Comment: Although normality of errors is not established we will proceed to evaluate the model performance

Model Validation: Holdout Method using RMSE

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
import numpy as np

rmse = (np.sqrt(mean_squared_error(boston_train['MEDV'], boston_train['pred'])))
r2 = r2_score(boston_train['MEDV'], boston_train['pred'])

print('RMSE is {}'.format(rmse))
```

```
RMSE is 4.828905476132975
```

```
print('R2 score is {}'.format(r2))
```

```
R2 score is 0.7315826585744917
```

```
y_test_predict = hp_model.predict(boston_test)
rmse = (np.sqrt(mean_squared_error(boston_test['MEDV'], y_test_predict)))
r2 = r2_score(boston_test['MEDV'], y_test_predict)
```

```
print('RMSE is {}'.format(rmse))
```

```
RMSE is 5.10498883392664
```

```
print('R2 score is {}'.format(r2))
```

```
R2 score is 0.6446261208488169
```

K-fold cross validation

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression

X = boston.drop(labels=['MEDV', 'ZN', 'INDUS', 'RAD', 'AGE'], axis=1)
Y = boston['MEDV']
hp_model1 = LinearRegression()

folds = KFold(n_splits = 4, shuffle = True, random_state = 100)
scores = cross_val_score(hp_model1, X, Y, scoring='r2', cv=folds)

print("Mean 4-Fold R Squared: {}".format(np.mean(scores)))
```

```
Mean 4-Fold R Squared: 0.7073913275172319
```

```
cv_rmse_scores= cross_val_score(hp_model1, X, Y, cv=folds, scoring='neg_mean_squared_error')
np.sqrt(-(np.mean(cv_rmse_scores)))
```

```
4.971825927756832
```

Comment: RMSE and R squared values using K-fold validation are similar to overall RMSE and R squared values

The model can be implemented for decision making